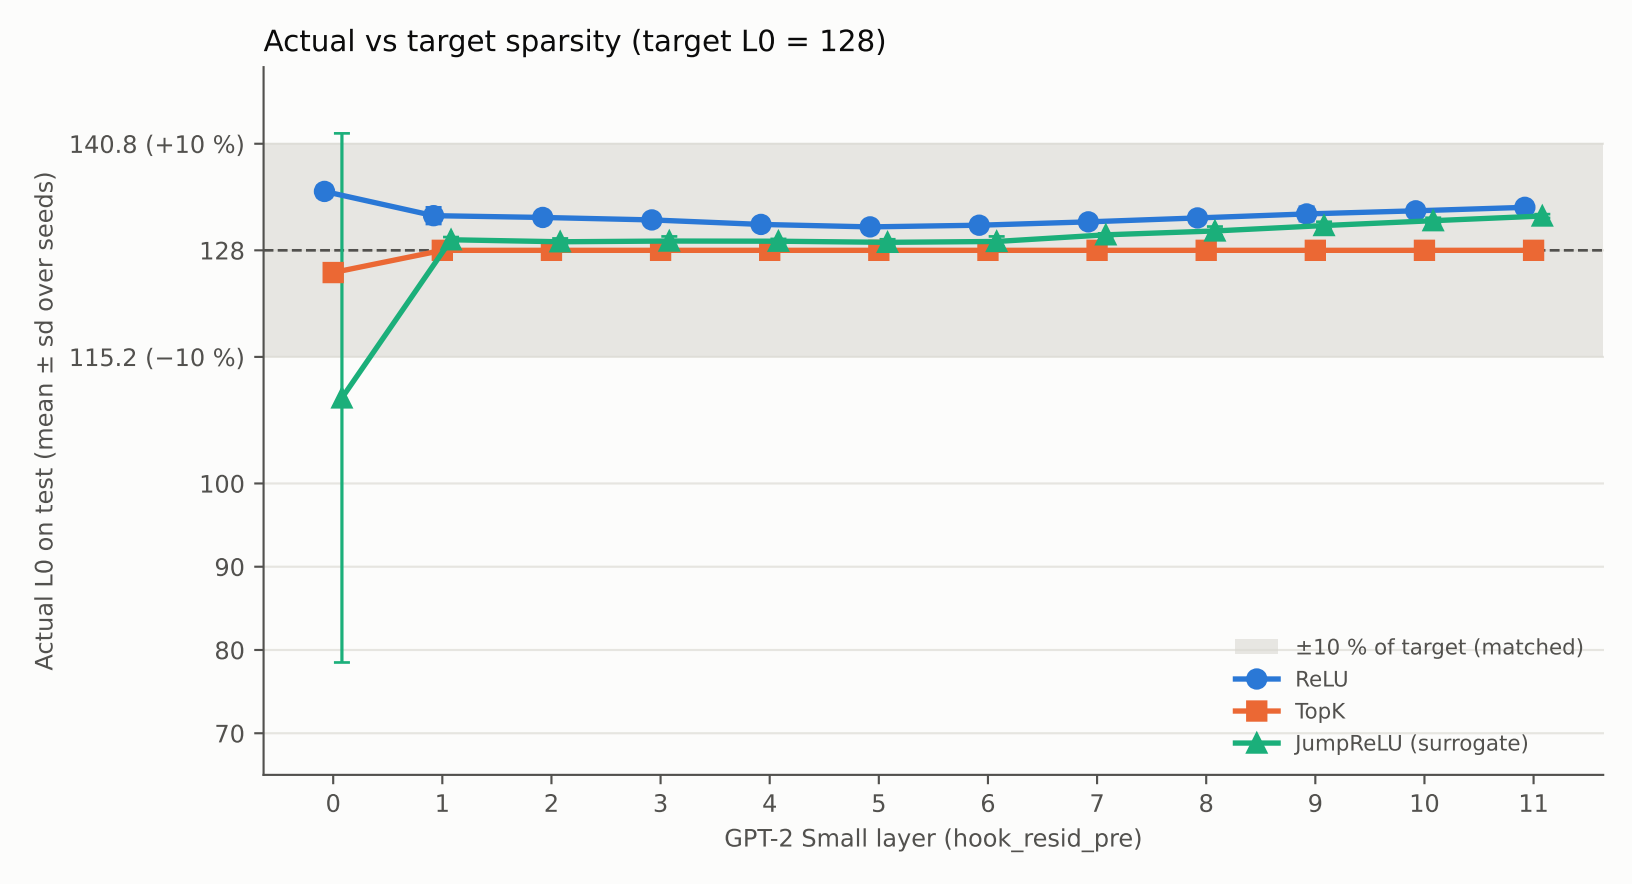
Actual vs target sparsity (target L0 = 128)
140.8 (+10 %)
128
115.2 (−10 %)
100
90
80
70
Actual L0 on test (mean ± sd over seeds)
0
1
2
3
4
5
6
7
8
9
10
11
GPT-2 Small layer (hook_resid_pre)
±10 % of target (matched)
ReLU
TopK
JumpReLU (surrogate)