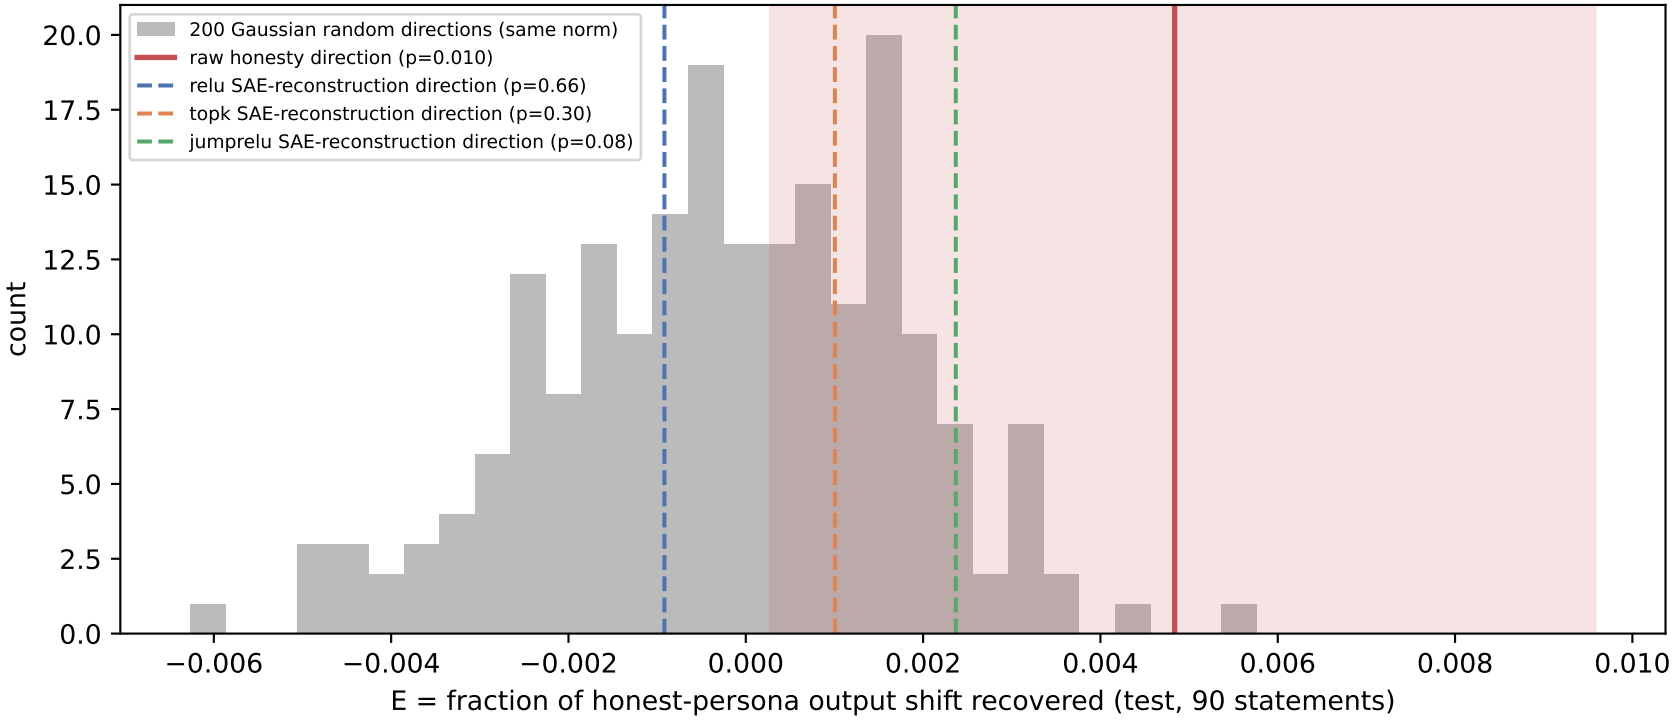
Task C — GPT-2 L6 final-token intervention, norm 0.22 (positive control: activation patch E=0.234, off-scale)
200 Gaussian random directions (same norm)
raw honesty direction (p=0.010)
relu SAE-reconstruction direction (p=0.66)
topk SAE-reconstruction direction (p=0.30)
jumprelu SAE-reconstruction direction (p=0.08)
count
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
20.0
−0.006
−0.004
−0.002
0.000
0.002
0.004
0.006
0.008
0.010
E = fraction of honest-persona output shift recovered (test, 90 statements)